

DSP3_Classification_final_dpmin4

August 26, 2020

```
[1]: import pandas as pd
import numpy as np
import sklearn
from sklearn.linear_model import LogisticRegressionCV
import seaborn as sns
%matplotlib inline
import matplotlib.pyplot as plt

[2]: data_links_biarc0 = ['https://dpass3ngrams.s3.amazonaws.com/csv/output_dpmin4.csv',
    'https://dpass3ngrams.s3.amazonaws.com/csv/output_dpmin5.csv',
    'https://dpass3ngrams.s3.amazonaws.com/csv/output_dpmin6.csv']

data_links_biarc1 = ['https://dpass3ngrams.s3.amazonaws.com/csv/output2_dpmin6.csv.
    ↳zip',
    'https://dpass3ngrams.s3.amazonaws.com/csv/output2_dpmin5.csv.zip',
    'https://dpass3ngrams.s3.amazonaws.com/csv/output2_dpmin4.csv.zip']

data_links = data_links_biarc1

[3]: # !wget https://dpass3ngrams.s3.amazonaws.com/csv/output2_dpmin6.csv.zip
# !wget https://dpass3ngrams.s3.amazonaws.com/csv/output2_dpmin5.csv.zip
!wget https://dpass3ngrams.s3.amazonaws.com/csv/output2_dpmin4.csv.zip

#TODO: generify this
```

--2020-08-21 14:30:12--

https://dpass3ngrams.s3.amazonaws.com/csv/output2_dpmin4.csv.zip

Resolving dpass3ngrams.s3.amazonaws.com (dpass3ngrams.s3.amazonaws.com)...

52.216.133.147

Connecting to dpass3ngrams.s3.amazonaws.com

(dpass3ngrams.s3.amazonaws.com)|52.216.133.147|:443... connected.

HTTP request sent, awaiting response... 200 OK

Length: 5705401 (5.4M) [application/zip]

Saving to: output2_dpmin4.csv.zip

output2_dpmin4.csv. 100%[=====>] 5.44M 3.58MB/s in 1.5s

2020-08-21 14:30:14 (3.58 MB/s) - output2_dpmin4.csv.zip saved
[5705401/5705401]

```
[4]: !unzip \*.zip  
  
!ls
```

Archive: output2_dpmin4.csv.zip
 inflating: output2_dpmin4.csv
DSP3_Classification.ipynb output2_dpmin4.csv
DSP3_Classification_final.ipynb output2_dpmin4.csv.zip

```
[5]: # _df = pd.read_csv ('output2_dpmin{0}.csv'.format(4))  
# print(df.head())
```

```
[6]: # words[4]['Word1']
```

We run this with 3 *DPMIN* Values: 4,5,6

```
[7]: checks = [4]
```

```
[8]: X = {}  
Y = {}  
words = {}  
embd_X = {}  
  
for i in checks:  
  
    df = pd.read_csv ('output2_dpmin{0}.csv'.format(i))  
    feature_len = len(df.columns)-3 #3 because we don't care about: Word1, Word2 and True_  
    →Label  
  
    X_df = df[['Feature_'+str(j) for j in range (feature_len)]]  
    Y_df = df[['True Label']]  
    words[i] = df[['Word1','Word2']]  
  
    X[i] = np.squeeze(X_df.to_numpy())  
    Y[i] = np.squeeze(Y_df.to_numpy())
```

0.1 Machine Learning

0.1.1 Preprocessing

```
[9]: from sklearn.decomposition import PCA  
  
K = 100 # the new dimesnion size - we reducing to  
  
for i in checks:
```

```
pca = PCA(n_components = K)
embd_X[i] = pca.fit_transform(X[i])
total_var = pca.explained_variance_ratio_.cumsum()[-1]
print('Total variance retrained with this PCA: {0}%'.format(total_var*100))
```


Total variance retrained with this PCA: 70.7462314695531%

Then, we will visualize the datapoints

```
[10]: from sklearn.manifold import TSNE

for i in checks:

    np.random.seed(33)
    tsne = TSNE(n_components=2).fit_transform(embd_X[i]) #TSNE is good for visualising
    →scatter plots
    plt.figure(figsize=(16,10))
    sns.scatterplot(tsne[:,0],tsne[:,1],hue=Y[i])
    plt.title('Scatter Plot for DPMIN={}'.format(i))
```



Doesn't seem like it can be easily seperated, at least with no kernel.

0.1.2 Training

We will train with logistic regression and with $CV = 10$

```
[11]: clf_logic = {}
      for i in checks:
          tmp_clf = LogisticRegressionCV(cv = 10,class_weight='balanced',max_iter=10_000).
          →fit(embd_X[i],Y[i])
          print('i-th score: ', tmp_clf.score(embd_X[i],Y[i]))
          clf_logic[i] = tmp_clf
```

i-th score: 0.6457220504170181

0.1.3 Analysis

```
[12]: from sklearn.metrics import precision_recall_fscore_support, confusion_matrix, roc_auc_score

      def analyse_clf (clf,i):


          y_pred = np.round(clf.predict(embd_X[i]))
          y_true = Y[i]

          prf = list(precision_recall_fscore_support( y_true, y_pred,average='binary'))[:-1]
          prf.append(roc_auc_score(y_true,y_pred))
          conf_mat = confusion_matrix(y_true, y_pred,normalize='true') #also normlize according_
          →to the truth

          print('Precision: {0}\nRecall: {1}\nF1_Score:{2}\n ROC: {3}'.format(*prf))
          print('-----')
          plt.figure()
          sns.heatmap(conf_mat,annot=True)
          # row is true, column is predicted

[13]: for i in checks:
          analyse_clf(clf_logic[i],i)
```

Precision: 0.867578751794612
 Recall: 0.6315627689659412
 F1_Score:0.7309922795033266
 ROC: 0.66132905981788



Now let us try *SVC*

```
[14]: scoring = ['precision', 'recall', 'accuracy', 'f1', 'roc_auc']
from sklearn.metrics import recall_score
from sklearn.metrics import precision_score
from sklearn.metrics import accuracy_score
from sklearn.metrics import roc_auc_score

from sklearn.model_selection import cross_validate
from sklearn.metrics import recall_score
from sklearn.metrics import precision_score
from sklearn.metrics import accuracy_score

from sklearn.svm import SVC
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import StandardScaler

best_models = {}

for i in checks:

    tmp_clf = make_pipeline(StandardScaler(), SVC(gamma='auto'))
    # tmp_clf.fit(embd_X[i], Y[i])

    scores = cross_validate(tmp_clf, embd_X[i], Y[i],
        ↳scoring=scoring, return_estimator=True, cv=10)
    print("*****{ }*****".format(i))
```

```
# print(scores)

best_model_ndx = np.argmax(scores['test_roc_auc'])
print
best_models[i] = scores['estimator'][best_model_ndx]

analyse_clf(best_models[i],i)
```


*****4*****

Precision: 0.7739326057821175

Recall: 0.9940366408459363

F1_Score:0.8702836535873837

ROC: 0.531789794024229



Let us try with *balanced* class weight

```
[15]: best_models_balanced = {}

for i in checks:

    tmp_clf = make_pipeline(StandardScaler(), SVC(gamma='auto',class_weight='balanced'))
    # tmp_clf.fit(embd_X[i],Y[i])

    scores = cross_validate(tmp_clf, embd_X[i], Y[i],
    ↳scoring=scoring,return_estimator=True,cv=10)
    print("*****{}*****".format(i))
```


```
# print(scores)

best_model_ndx = np.argmax(scores['test_roc_auc'])
print
best_models_balanced[i] = scores['estimator'][best_model_ndx]

analyse_clf(best_models_balanced[i],i)

best_clf = best_models_balanced[4]
```

```
*****4*****
Precision: 0.8787528868360277
Recall: 0.6549858600762326
F1_Score:0.75054596688975
ROC: 0.6826938756645938
-----
```



```
[16]: from collections import Counter

for i in checks:
    print("*****{}*****".format(i))

    counted_vals = Counter(Y[i])
    zero_label_percent = counted_vals[0] / (counted_vals[0] + counted_vals[1])
    one_label_percent = counted_vals[1] / (counted_vals[0] + counted_vals[1])
    print("No-hypernyms",zero_label_percent)
```

```
print("hypernyms",one_label_percent)
```

```
*****4*****
```

```
No-hypernyms 0.2378408771436604
```

```
hypernyms 0.7621591228563396
```

```
[17]: # We supposed to play with the examples now  
embd_X
```

```
[17]: {4: array([[ -0.10029389, -0.13539825, -0.384451 , ...,  0.00271945,  
                0.0013594 ,  0.00175056],  
               [ -0.19144404,  0.10003238, -0.28997174, ...,  0.00328079,  
                -0.00307037,  0.00269231],  
               [ -0.29098437, -0.05802356,  0.19640241, ...,  0.04034176,  
                0.04514644, -0.02322879],  
               ...,  
               [ -0.13210339, -0.15700779, -0.32373864, ...,  0.03877869,  
                0.01598212,  0.01070661],  
               [ -0.40883516, -0.59908823,  0.2890192 , ...,  0.00262066,  
                0.0015617 , -0.00371533],  
               [ -0.36605547,  0.68059671, -0.0910479 , ...,  0.00324655,  
                0.00232376, -0.00320101]])}
```

```
[18]: # extract Bad and Good Examples:
```

```
TP, FP, TN, FN = [], [], [], []
```

```
got_enough_examples = False
```

```
predictions = best_clf.predict(embd_X[4])
```

```
counter = 0
```

```
for w1,w2,pred,true in zip(words[4]['Word1'],words[4]['Word2'],predictions,Y[4]):
```

```
    if true == 0:
```

```
        if pred == 0:
```

```
            TN.append((w1,w2))
```

```
        if pred == 1:
```

```
            FP.append((w1,w2))
```

```
    if true == 1:
```

```
        if pred == 0:
```

```
            FN.append((w1,w2))
```

```
        if pred == 1:
```

```
            TP.append((w1,w2))
```

```
[18]:
```